

# week\_5\_assignment 5

April 17, 2022

## 1 Week5: Assignments

### 1.0.1 Author: Ganesh Kale

import required packages

```
[22]: from tensorflow import keras
      from keras import models
      from keras import layers
      from keras.datasets import imdb
      import numpy as np
      import matplotlib.pyplot as plt
      from tensorflow.keras.utils import to_categorical

      from IPython.core.interactiveshell import InteractiveShell
      InteractiveShell.ast_node_interactivity = 'all'
```

## 2 Assignment 5.1: Classifying movie reviews: A binary classification example

```
[4]: # load the data

      (train_data, train_labels), (test_data, test_labels) = imdb.
      ↪load_data(num_words=10000)
```

```
[6]: # display train and test data sample

      print(train_data[0])
      train_labels[0]
```

```
[1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 36,
256, 5, 25, 100, 43, 838, 112, 50, 670, 2, 9, 35, 480, 284, 5, 150, 4, 172, 112,
167, 2, 336, 385, 39, 4, 172, 4536, 1111, 17, 546, 38, 13, 447, 4, 192, 50, 16,
6, 147, 2025, 19, 14, 22, 4, 1920, 4613, 469, 4, 22, 71, 87, 12, 16, 43, 530,
38, 76, 15, 13, 1247, 4, 22, 17, 515, 17, 12, 16, 626, 18, 2, 5, 62, 386, 12, 8,
316, 8, 106, 5, 4, 2223, 5244, 16, 480, 66, 3785, 33, 4, 130, 12, 16, 38, 619,
5, 25, 124, 51, 36, 135, 48, 25, 1415, 33, 6, 22, 12, 215, 28, 77, 52, 5, 14,
```

```
407, 16, 82, 2, 8, 4, 107, 117, 5952, 15, 256, 4, 2, 7, 3766, 5, 723, 36, 71,
43, 530, 476, 26, 400, 317, 46, 7, 4, 2, 1029, 13, 104, 88, 4, 381, 15, 297, 98,
32, 2071, 56, 26, 141, 6, 194, 7486, 18, 4, 226, 22, 21, 134, 476, 26, 480, 5,
144, 30, 5535, 18, 51, 36, 28, 224, 92, 25, 104, 4, 226, 65, 16, 38, 1334, 88,
12, 16, 283, 5, 16, 4472, 113, 103, 32, 15, 16, 5345, 19, 178, 32]
```

```
[6]: 1
```

### 2.0.1 Encoding the integer sequences into a binary matrix

```
[7]: # create function to encode sequence

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
[8]: # display sample of training data after encoding

x_train[0]
```

```
[8]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
[9]: # vectorize the labels

y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

### 2.0.2 Build the NN Model

```
[12]: model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

#### Compile the model

```
[13]: model.compile(optimizer="rmsprop",
                    loss="binary_crossentropy",
                    metrics=["accuracy"])
```

#### Splitting data into training and validation

```
[14]: x_val = x_train[:10000]
      partial_x_train = x_train[10000:]
      y_val = y_train[:10000]
      partial_y_train = y_train[10000:]
```

### Train the model

```
[15]: history = model.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))
```

```
Epoch 1/20
30/30 [=====] - 1s 45ms/step - loss: 0.5145 - accuracy:
0.7843 - val_loss: 0.4133 - val_accuracy: 0.8416
Epoch 2/20
30/30 [=====] - 1s 29ms/step - loss: 0.3061 - accuracy:
0.8999 - val_loss: 0.3046 - val_accuracy: 0.8861
Epoch 3/20
30/30 [=====] - 1s 27ms/step - loss: 0.2234 - accuracy:
0.9269 - val_loss: 0.2809 - val_accuracy: 0.8907
Epoch 4/20
30/30 [=====] - 1s 29ms/step - loss: 0.1788 - accuracy:
0.9412 - val_loss: 0.2736 - val_accuracy: 0.8909
Epoch 5/20
30/30 [=====] - 1s 29ms/step - loss: 0.1431 - accuracy:
0.9551 - val_loss: 0.2816 - val_accuracy: 0.8861
Epoch 6/20
30/30 [=====] - 1s 38ms/step - loss: 0.1169 - accuracy:
0.9638 - val_loss: 0.3234 - val_accuracy: 0.8787
Epoch 7/20
30/30 [=====] - 1s 42ms/step - loss: 0.0989 - accuracy:
0.9711 - val_loss: 0.3150 - val_accuracy: 0.8816
Epoch 8/20
30/30 [=====] - 1s 32ms/step - loss: 0.0814 - accuracy:
0.9767 - val_loss: 0.3298 - val_accuracy: 0.8805
Epoch 9/20
30/30 [=====] - 1s 30ms/step - loss: 0.0659 - accuracy:
0.9834 - val_loss: 0.3539 - val_accuracy: 0.8804
Epoch 10/20
30/30 [=====] - 1s 30ms/step - loss: 0.0552 - accuracy:
0.9863 - val_loss: 0.3963 - val_accuracy: 0.8706
Epoch 11/20
30/30 [=====] - 1s 33ms/step - loss: 0.0447 - accuracy:
0.9897 - val_loss: 0.4044 - val_accuracy: 0.8756
Epoch 12/20
30/30 [=====] - 1s 31ms/step - loss: 0.0358 - accuracy:
```

```

0.9927 - val_loss: 0.4500 - val_accuracy: 0.8685
Epoch 13/20
30/30 [=====] - 1s 28ms/step - loss: 0.0282 - accuracy:
0.9947 - val_loss: 0.4663 - val_accuracy: 0.8726
Epoch 14/20
30/30 [=====] - 1s 27ms/step - loss: 0.0231 - accuracy:
0.9957 - val_loss: 0.5182 - val_accuracy: 0.8726
Epoch 15/20
30/30 [=====] - 1s 28ms/step - loss: 0.0218 - accuracy:
0.9950 - val_loss: 0.5320 - val_accuracy: 0.8696
Epoch 16/20
30/30 [=====] - 1s 29ms/step - loss: 0.0130 - accuracy:
0.9984 - val_loss: 0.5733 - val_accuracy: 0.8700
Epoch 17/20
30/30 [=====] - 1s 34ms/step - loss: 0.0124 - accuracy:
0.9981 - val_loss: 0.5943 - val_accuracy: 0.8674
Epoch 18/20
30/30 [=====] - 1s 30ms/step - loss: 0.0068 - accuracy:
0.9997 - val_loss: 0.6573 - val_accuracy: 0.8677
Epoch 19/20
30/30 [=====] - 1s 33ms/step - loss: 0.0114 - accuracy:
0.9971 - val_loss: 0.6667 - val_accuracy: 0.8669
Epoch 20/20
30/30 [=====] - 1s 27ms/step - loss: 0.0038 - accuracy:
0.9999 - val_loss: 0.6906 - val_accuracy: 0.8645

```

```

[16]: # see the training history

history_dict = history.history
history_dict.keys()

```

```

[16]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```

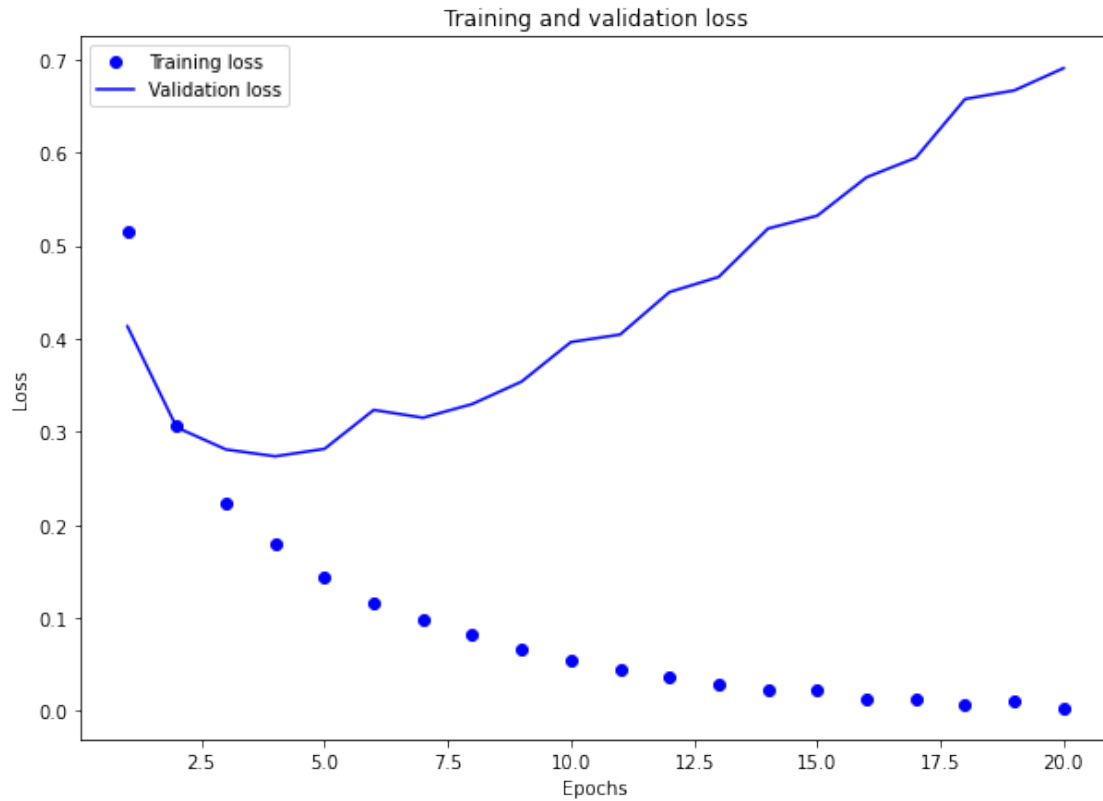
### Plotting the training and validation loss

```

[20]: plt.figure(figsize=(10,7))

history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show();

```



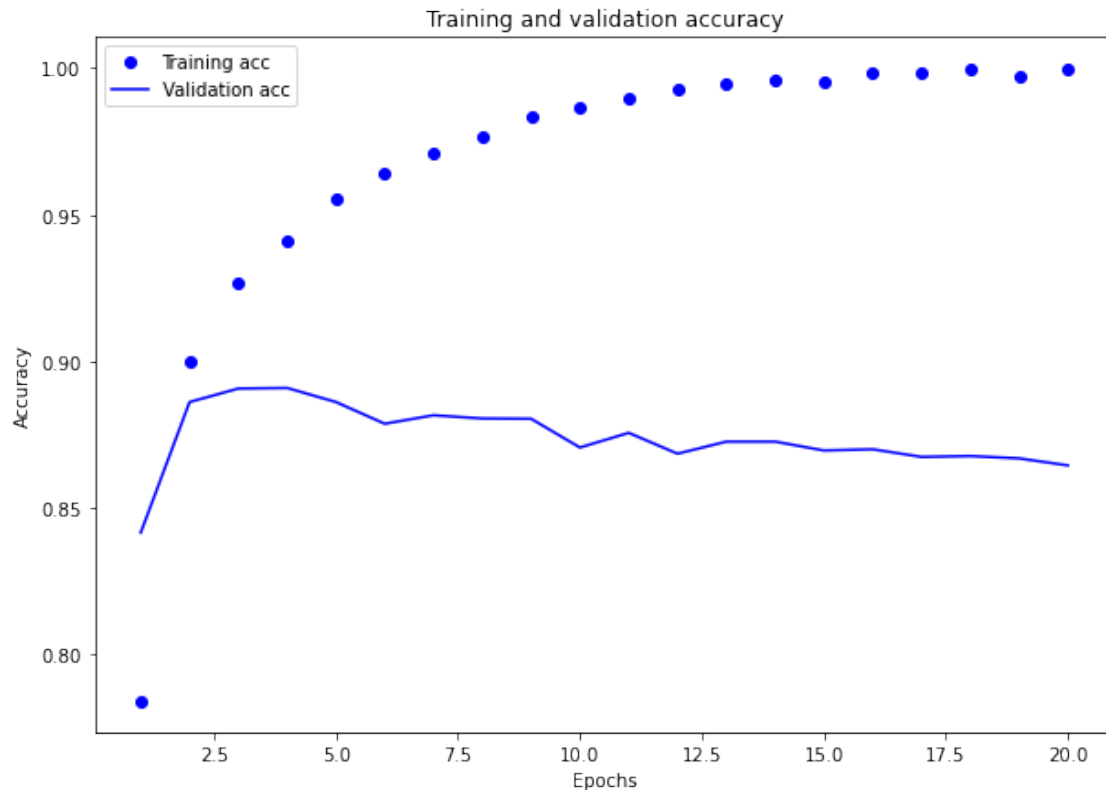
### Plotting the training and validation accuracy

```
[22]: plt.clf()

plt.figure(figsize=(10,7))

acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show();
```

<Figure size 432x288 with 0 Axes>



### Retraining a model from scratch - for 4 epochs

```
[23]: model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

Epoch 1/4

49/49 [=====] - 0s 10ms/step - loss: 0.5049 - accuracy: 0.7966

Epoch 2/4

49/49 [=====] - 0s 8ms/step - loss: 0.2975 - accuracy: 0.9072

Epoch 3/4

49/49 [=====] - 0s 9ms/step - loss: 0.2170 - accuracy: 0.9286

```
Epoch 4/4
49/49 [=====] - 0s 8ms/step - loss: 0.1783 - accuracy:
0.9394
```

```
[23]: <tensorflow.python.keras.callbacks.History at 0x7f2911e17880>
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.3007 -
accuracy: 0.8813
```

```
[25]: # display the result
```

```
results
```

```
[25]: [0.3007475733757019, 0.8813199996948242]
```

```
[26]: print(f"The accuracy of the model is {round(results[1],2)*100}%")
```

```
The accuracy of the model is 88.0%
```

### 2.0.3 END

## 3 Assignment 5.2: Classifying newswires: A multiclass classification example

```
[8]: # loading the Reuters dataset
```

```
from tensorflow.keras.datasets import reuters
(train_data, train_labels), (test_data, test_labels) = reuters.load_data(
    num_words=10000)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/reuters.npz
2113536/2110848 [=====] - 0s 0us/step
```

```
[9]: # display records in each data set train and test
```

```
len(train_data)
len(test_data)
```

```
[9]: 8982
```

```
[9]: 2246
```

```
[11]: # display train data sample
```

```
print(train_data[10])
```

```
[1, 245, 273, 207, 156, 53, 74, 160, 26, 14, 46, 296, 26, 39, 74, 2979, 3554,
14, 46, 4689, 4329, 86, 61, 3499, 4795, 14, 61, 451, 4329, 17, 12]
```

```
[12]: # Decoding newswires back to text

word_index = reuters.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
decoded_newswire = " ".join([reverse_word_index.get(i - 3, "?") for i in
    ↪train_data[0]])

decoded_newswire
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/reuters\\_word\\_index.json](https://storage.googleapis.com/tensorflow/tf-keras-datasets/reuters_word_index.json)  
557056/550378 [=====] - 0s 1us/step

```
[12]: ' ? ? said as a result of its december acquisition of space co it expects
earnings per share in 1987 of 1 15 to 1 30 dlrs per share up from 70 cts in 1986
the company said pretax net should rise to nine to 10 mln dlrs from six mln dlrs
in 1986 and rental operation revenues to 19 to 22 mln dlrs from 12 5 mln dlrs it
said cash flow per share this year should be 2 50 to three dlrs reuter 3'
```

```
[13]: # label associated with an example is an integer between 0 and 45

train_labels[10]
```

```
[13]: 3
```

Encoding the input data

```
[18]: def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1
    return results
```

```
[19]: x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

Encoding the labels

```
[20]: def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results
y_train = to_one_hot(train_labels)
y_test = to_one_hot(test_labels)
```



```
[23]: y_train = to_categorical(train_labels)
      y_test = to_categorical(test_labels)
```

Building your model

```
[25]: model = keras.Sequential([
      layers.Dense(64, activation="relu"),
      layers.Dense(64, activation="relu"),
      layers.Dense(46, activation="softmax")
    ])
```

Compiling the model

```
[26]: model.compile(optimizer="rmsprop",
      loss="categorical_crossentropy",
      metrics=["accuracy"])
```

Setting aside a validation set

```
[27]: x_val = x_train[:1000]
      partial_x_train = x_train[1000:]
      y_val = y_train[:1000]
      partial_y_train = y_train[1000:]
```

```
[28]: history = model.fit(partial_x_train,
      partial_y_train,
      epochs=20,
      batch_size=512,
      validation_data=(x_val, y_val))
```

Epoch 1/20

16/16 [=====] - 1s 35ms/step - loss: 2.6546 - accuracy: 0.5004 - val\_loss: 1.7578 - val\_accuracy: 0.6500

Epoch 2/20

16/16 [=====] - 1s 44ms/step - loss: 1.4382 - accuracy: 0.7015 - val\_loss: 1.3042 - val\_accuracy: 0.7170

Epoch 3/20

16/16 [=====] - 1s 45ms/step - loss: 1.0600 - accuracy: 0.7776 - val\_loss: 1.1244 - val\_accuracy: 0.7550

Epoch 4/20

16/16 [=====] - 1s 44ms/step - loss: 0.8344 - accuracy: 0.8217 - val\_loss: 1.0341 - val\_accuracy: 0.7860

Epoch 5/20

16/16 [=====] - 0s 20ms/step - loss: 0.6638 - accuracy: 0.8602 - val\_loss: 0.9458 - val\_accuracy: 0.7990

Epoch 6/20

16/16 [=====] - 0s 19ms/step - loss: 0.5292 - accuracy: 0.8910 - val\_loss: 0.8961 - val\_accuracy: 0.8110

```

Epoch 7/20
16/16 [=====] - 0s 18ms/step - loss: 0.4262 - accuracy:
0.9137 - val_loss: 0.8996 - val_accuracy: 0.8090
Epoch 8/20
16/16 [=====] - 0s 19ms/step - loss: 0.3465 - accuracy:
0.9286 - val_loss: 0.8951 - val_accuracy: 0.8090
Epoch 9/20
16/16 [=====] - 0s 17ms/step - loss: 0.2855 - accuracy:
0.9382 - val_loss: 0.8836 - val_accuracy: 0.8210
Epoch 10/20
16/16 [=====] - 0s 16ms/step - loss: 0.2423 - accuracy:
0.9451 - val_loss: 0.8826 - val_accuracy: 0.8230
Epoch 11/20
16/16 [=====] - 0s 16ms/step - loss: 0.2070 - accuracy:
0.9506 - val_loss: 0.9029 - val_accuracy: 0.8210
Epoch 12/20
16/16 [=====] - 0s 15ms/step - loss: 0.1823 - accuracy:
0.9521 - val_loss: 0.9517 - val_accuracy: 0.8090
Epoch 13/20
16/16 [=====] - 0s 18ms/step - loss: 0.1664 - accuracy:
0.9533 - val_loss: 0.9173 - val_accuracy: 0.8130
Epoch 14/20
16/16 [=====] - 1s 34ms/step - loss: 0.1526 - accuracy:
0.9546 - val_loss: 1.0090 - val_accuracy: 0.7960
Epoch 15/20
16/16 [=====] - 1s 39ms/step - loss: 0.1366 - accuracy:
0.9560 - val_loss: 0.9653 - val_accuracy: 0.8130
Epoch 16/20
16/16 [=====] - 1s 50ms/step - loss: 0.1333 - accuracy:
0.9569 - val_loss: 0.9726 - val_accuracy: 0.8190
Epoch 17/20
16/16 [=====] - 1s 34ms/step - loss: 0.1273 - accuracy:
0.9557 - val_loss: 1.0120 - val_accuracy: 0.8000
Epoch 18/20
16/16 [=====] - 0s 26ms/step - loss: 0.1205 - accuracy:
0.9567 - val_loss: 1.0046 - val_accuracy: 0.8150
Epoch 19/20
16/16 [=====] - 0s 23ms/step - loss: 0.1172 - accuracy:
0.9572 - val_loss: 1.1131 - val_accuracy: 0.7870
Epoch 20/20
16/16 [=====] - 0s 22ms/step - loss: 0.1106 - accuracy:
0.9563 - val_loss: 1.0359 - val_accuracy: 0.8060

```

Plotting the training and validation loss

```

[30]: plt.figure(figsize=(10,7))


      loss = history.history["loss"]

```

```

val_loss = history.history["val_loss"]
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show();

```



Plotting the training and validation accuracy

```

[31]: plt.clf()


plt.figure(figsize=(10,7))

acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")

```

```
plt.ylabel("Accuracy")
plt.legend()
plt.show();
```

<Figure size 432x288 with 0 Axes>



Retraining a model from scratch

```
[32]: model = keras.Sequential([
    layers.Dense(64, activation="relu"),
    layers.Dense(64, activation="relu"),
    layers.Dense(46, activation="softmax")
])
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics=["accuracy"])
model.fit(x_train,
          y_train,
          epochs=9,
          batch_size=512)
results = model.evaluate(x_test, y_test)
```

```

Epoch 1/9
18/18 [=====] - 0s 16ms/step - loss: 2.6042 - accuracy:
0.4963
Epoch 2/9
18/18 [=====] - 0s 13ms/step - loss: 1.4033 - accuracy:
0.7088
Epoch 3/9
18/18 [=====] - 0s 12ms/step - loss: 1.0273 - accuracy:
0.7808
Epoch 4/9
18/18 [=====] - 0s 13ms/step - loss: 0.7990 - accuracy:
0.8292
Epoch 5/9
18/18 [=====] - 0s 12ms/step - loss: 0.6345 - accuracy:
0.8661
Epoch 6/9
18/18 [=====] - 0s 13ms/step - loss: 0.5035 - accuracy:
0.8965
Epoch 7/9
18/18 [=====] - 0s 13ms/step - loss: 0.4030 - accuracy:
0.9144
Epoch 8/9
18/18 [=====] - 0s 13ms/step - loss: 0.3334 - accuracy:
0.9259
Epoch 9/9
18/18 [=====] - 0s 15ms/step - loss: 0.2758 - accuracy:
0.9378

```

[32]: <tensorflow.python.keras.callbacks.History at 0x7fcde67db550>

```

71/71 [=====] - 0s 2ms/step - loss: 0.9557 - accuracy:
0.7907

```

[33]: *# display the result*

```
results
```

[33]: [0.9557026028633118, 0.790739119052887]

[34]: `print(f"The accuracy of the model is {round(results[1],2)*100}%")`

The accuracy of the model is 79.0%

predictions on new data

[35]: *# predictions is a vector of length 46*

```
predictions = model.predict(x_test)
```

```
predictions[0].shape
```

[35]: (46,)

```
[36]: # The coefficients in this vector sum to 1  
  
np.sum(predictions[0])
```

[36]: 0.99999994

```
[37]: # the class with the highest probability  
  
np.argmax(predictions[0])
```

[37]: 3

different way to handle the labels and the loss

```
[38]: y_train = np.array(train_labels)  
y_test = np.array(test_labels)
```

```
[39]: model.compile(optimizer="rmsprop",  
                    loss="sparse_categorical_crossentropy",  
                    metrics=["accuracy"])
```

importance of having sufficiently large intermediate layers

```
[40]: model = keras.Sequential([  
    layers.Dense(64, activation="relu"),  
    layers.Dense(4, activation="relu"),  
    layers.Dense(46, activation="softmax")  
)  
model.compile(optimizer="rmsprop",  
              loss="categorical_crossentropy",  
              metrics=["accuracy"])  
model.fit(partial_x_train,  
          partial_y_train,  
          epochs=20,  
          batch_size=128,  
          validation_data=(x_val, y_val))
```

Epoch 1/20

63/63 [=====] - 1s 9ms/step - loss: 3.0746 - accuracy:  
0.2279 - val\_loss: 2.5696 - val\_accuracy: 0.2740

Epoch 2/20

63/63 [=====] - 1s 8ms/step - loss: 1.9711 - accuracy:  
0.5382 - val\_loss: 1.6253 - val\_accuracy: 0.6230

Epoch 3/20

63/63 [=====] - 1s 11ms/step - loss: 1.4587 - accuracy: 0.6419 - val\_loss: 1.4765 - val\_accuracy: 0.6360  
Epoch 4/20  
63/63 [=====] - 1s 9ms/step - loss: 1.2913 - accuracy: 0.6704 - val\_loss: 1.4255 - val\_accuracy: 0.6600  
Epoch 5/20  
63/63 [=====] - 1s 8ms/step - loss: 1.1830 - accuracy: 0.7005 - val\_loss: 1.4011 - val\_accuracy: 0.6620  
Epoch 6/20  
63/63 [=====] - 0s 7ms/step - loss: 1.0987 - accuracy: 0.7076 - val\_loss: 1.3835 - val\_accuracy: 0.6610  
Epoch 7/20  
63/63 [=====] - 1s 9ms/step - loss: 1.0299 - accuracy: 0.7103 - val\_loss: 1.4036 - val\_accuracy: 0.6570  
Epoch 8/20  
63/63 [=====] - 1s 8ms/step - loss: 0.9731 - accuracy: 0.7210 - val\_loss: 1.4025 - val\_accuracy: 0.6680  
Epoch 9/20  
63/63 [=====] - 0s 8ms/step - loss: 0.9242 - accuracy: 0.7388 - val\_loss: 1.4112 - val\_accuracy: 0.6780  
Epoch 10/20  
63/63 [=====] - 1s 8ms/step - loss: 0.8834 - accuracy: 0.7523 - val\_loss: 1.4567 - val\_accuracy: 0.6750  
Epoch 11/20  
63/63 [=====] - 1s 9ms/step - loss: 0.8461 - accuracy: 0.7602 - val\_loss: 1.4977 - val\_accuracy: 0.6790  
Epoch 12/20  
63/63 [=====] - 1s 8ms/step - loss: 0.8124 - accuracy: 0.7669 - val\_loss: 1.5203 - val\_accuracy: 0.6800  
Epoch 13/20  
63/63 [=====] - 0s 7ms/step - loss: 0.7838 - accuracy: 0.7750 - val\_loss: 1.5645 - val\_accuracy: 0.6700  
Epoch 14/20  
63/63 [=====] - 1s 8ms/step - loss: 0.7546 - accuracy: 0.7830 - val\_loss: 1.6169 - val\_accuracy: 0.6750  
Epoch 15/20  
63/63 [=====] - 0s 7ms/step - loss: 0.7325 - accuracy: 0.7908 - val\_loss: 1.6323 - val\_accuracy: 0.6690  
Epoch 16/20  
63/63 [=====] - 0s 8ms/step - loss: 0.7074 - accuracy: 0.7984 - val\_loss: 1.6979 - val\_accuracy: 0.6720  
Epoch 17/20  
63/63 [=====] - 1s 8ms/step - loss: 0.6895 - accuracy: 0.8029 - val\_loss: 1.7345 - val\_accuracy: 0.6670  
Epoch 18/20  
63/63 [=====] - 0s 7ms/step - loss: 0.6701 - accuracy: 0.8112 - val\_loss: 1.7949 - val\_accuracy: 0.6680  
Epoch 19/20

```
63/63 [=====] - 0s 7ms/step - loss: 0.6530 - accuracy:
0.8142 - val_loss: 1.8719 - val_accuracy: 0.6690
Epoch 20/20
63/63 [=====] - 0s 7ms/step - loss: 0.6363 - accuracy:
0.8167 - val_loss: 1.8490 - val_accuracy: 0.6830
```

```
[40]: <tensorflow.python.keras.callbacks.History at 0x7fcde76fa970>
```

## 4 END

## 5 Assignment 5.3: Predicting house prices: A regression example

```
[41]: # Loading the Boston housing dataset
```

```
from tensorflow.keras.datasets import boston_housing
(train_data, train_targets), (test_data, test_targets) = boston_housing.
    ↪load_data()
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston\\_housing.npz](https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston_housing.npz)

```
57344/57026 [=====] - 0s 1us/step
```

```
[42]: # display shape of dataset
```

```
train_data.shape
test_data.shape
```

```
[42]: (404, 13)
```

```
[42]: (102, 13)
```

```
[44]: # display targets - sample
```

```
train_targets[:5]
```

```
[44]: array([15.2, 42.3, 50. , 21.1, 17.7])
```

Normalizing the data

```
[45]: mean = train_data.mean(axis=0)
train_data -= mean
std = train_data.std(axis=0)
train_data /= std
test_data -= mean
test_data /= std
```



Building the model

```
[46]: def build_model():
      model = keras.Sequential([
          layers.Dense(64, activation="relu"),
          layers.Dense(64, activation="relu"),
          layers.Dense(1)
      ])
      model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
      return model
```

Validating approach using K-fold validation

```
[47]: k = 4
      num_val_samples = len(train_data) // k
      num_epochs = 100
      all_scores = []
      for i in range(k):
          print(f"Processing fold #{i}")
          val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
          val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
          partial_train_data = np.concatenate(
              [train_data[:i * num_val_samples],
               train_data[(i + 1) * num_val_samples:]],
              axis=0)
          partial_train_targets = np.concatenate(
              [train_targets[:i * num_val_samples],
               train_targets[(i + 1) * num_val_samples:]],
              axis=0)
          model = build_model()
          model.fit(partial_train_data, partial_train_targets,
                    epochs=num_epochs, batch_size=16, verbose=0)
          val_mse, val_mae = model.evaluate(val_data, val_targets, verbose=0)
          all_scores.append(val_mae)
```

Processing fold #0

```
[47]: <tensorflow.python.keras.callbacks.History at 0x7fcde5fdbcd0>
```

Processing fold #1

```
[47]: <tensorflow.python.keras.callbacks.History at 0x7fcde62af0a0>
```

Processing fold #2

```
[47]: <tensorflow.python.keras.callbacks.History at 0x7fcde4ba7ac0>
```

Processing fold #3

```
[47]: <tensorflow.python.keras.callbacks.History at 0x7fcde4a73880>
```

```
[48]: all_scores
```

```
[48]: [1.9127342700958252,  
      2.2901790142059326,  
      2.5055272579193115,  
      2.2678756713867188]
```

```
[49]: np.mean(all_scores)
```

```
[49]: 2.244079053401947
```

Saving the validation logs at each fold

```
[50]: num_epochs = 500  
all_mae_histories = []  
for i in range(k):  
    print(f"Processing fold #{i}")  
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]  
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]  
    partial_train_data = np.concatenate(  
        [train_data[:i * num_val_samples],  
         train_data[(i + 1) * num_val_samples:]],  
        axis=0)  
    partial_train_targets = np.concatenate(  
        [train_targets[:i * num_val_samples],  
         train_targets[(i + 1) * num_val_samples:]],  
        axis=0)  
    model = build_model()  
    history = model.fit(partial_train_data, partial_train_targets,  
                        validation_data=(val_data, val_targets),  
                        epochs=num_epochs, batch_size=16, verbose=0)  
    mae_history = history.history["val_mae"]  
    all_mae_histories.append(mae_history)
```

Processing fold #0

Processing fold #1

Processing fold #2

Processing fold #3


Building the history of successive mean K-fold validation scores

```
[51]: average_mae_history = [  
    np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)]  
p
```

Plotting validation scores

```
[54]: plt.figure(figsize=(10,7))


plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)
plt.xlabel("Epochs")
plt.ylabel("Validation MAE")
plt.show();
```



Plotting validation scores, excluding the first 10 data points


```
[56]: truncated_mae_history = average_mae_history[10:]
plt.figure(figsize=(10,7))

plt.plot(range(1, len(truncated_mae_history) + 1), truncated_mae_history)
plt.xlabel("Epochs")
plt.ylabel("Validation MAE")
plt.show();
```



```
[62]: def smooth_curve(points, factor=0.9):  
    smoothed_points = []  
    for point in points:  
        if smoothed_points:  
            previous = smoothed_points[-1]  
            smoothed_points.append(previous * factor + point * (1 - factor))  
        else:  
            smoothed_points.append(point)  
    return smoothed_points
```

```
[64]: smooth_mae_history = smooth_curve(average_mae_history[10:])  
plt.figure(figsize=(10,7))  
  
plt.plot(range(1, len(smooth_mae_history) + 1), smooth_mae_history)  
plt.xlabel('Epochs')  
plt.ylabel('Validation MAE')  
plt.show();
```



Training the final model

```
[57]: model = build_model()
      model.fit(train_data, train_targets,
                epochs=130, batch_size=16, verbose=0)
      test_mse_score, test_mae_score = model.evaluate(test_data, test_targets)
```

```
[57]: <tensorflow.python.keras.callbacks.History at 0x7fcdefa8dd00>
```

```
4/4 [=====] - 0s 1ms/step - loss: 13.5448 - mae: 2.4619
```

```
[58]: test_mae_score
```

```
[58]: 2.461940288543701
```

predictions on new data

```
[59]: predictions = model.predict(test_data)
      predictions[0]
```

```
[59]: array([7.760147], dtype=float32)
```

**6   END**